



Empowering .NET Applications with Semantic Kernel and Azure OpenAI

Orestis Meikopoulos

Head of Engineering @ Code Create

<https://linkedin.com/in/ormikopo>



Agenda

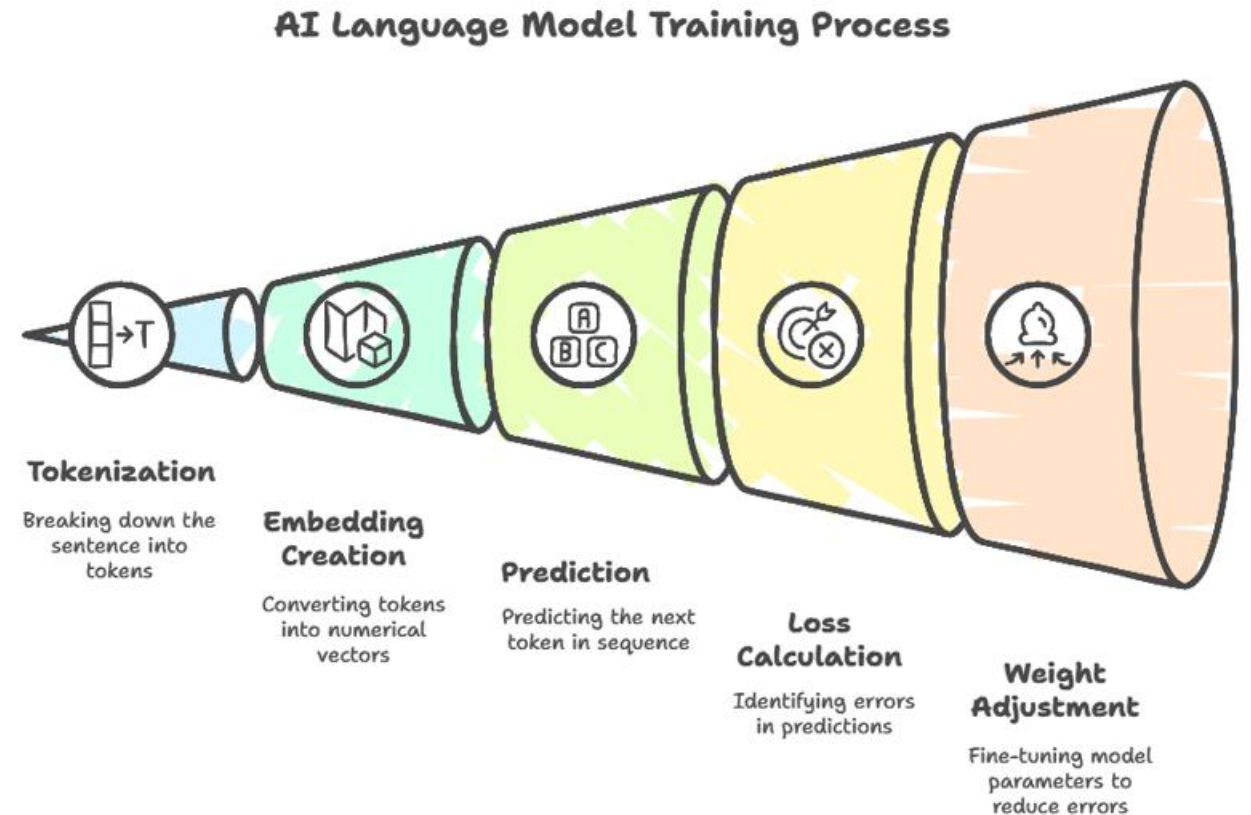
- Key concepts:
 - Generative AI & LLMs
 - Tokens
 - Embeddings
 - Vector databases
 - Retrieval Augmented Generation
 - OpenAI function calling
- Semantic Kernel
- Azure AI Agent Service
- Demo

How Generative AI & LLMs Work

- **Generative AI in a nutshell:** Creates new text, code, images or audio from a user prompt instead of just classifying or searching.
- **Everything starts with a model:** Trained on huge datasets, spots patterns and produces statistically similar but fresh output.
- **LLMs - The language engine:** A Generative AI app uses an LLM for natural-language I/O; other specialist models handle images or audio.

How Generative AI & LLMs Work

- **Under the hood - The token loop:**
 - Prompt is broken into tokens (word pieces).
 - Each token gets an embedding (aka, a high-dimensional meaning vector).
 - The model predicts the next token, one step at a time, updating like super-powered autocomplete.
- **Learning = minimizing “oops”:**
During training the model sees the real next token, measures the error (loss) and nudges its weights to do better (millions of times over).
- **What you can build:** Natural-language generation, image, audio, and code creation (e.g., ChatGPT writes a C# tic-tac-toe game).

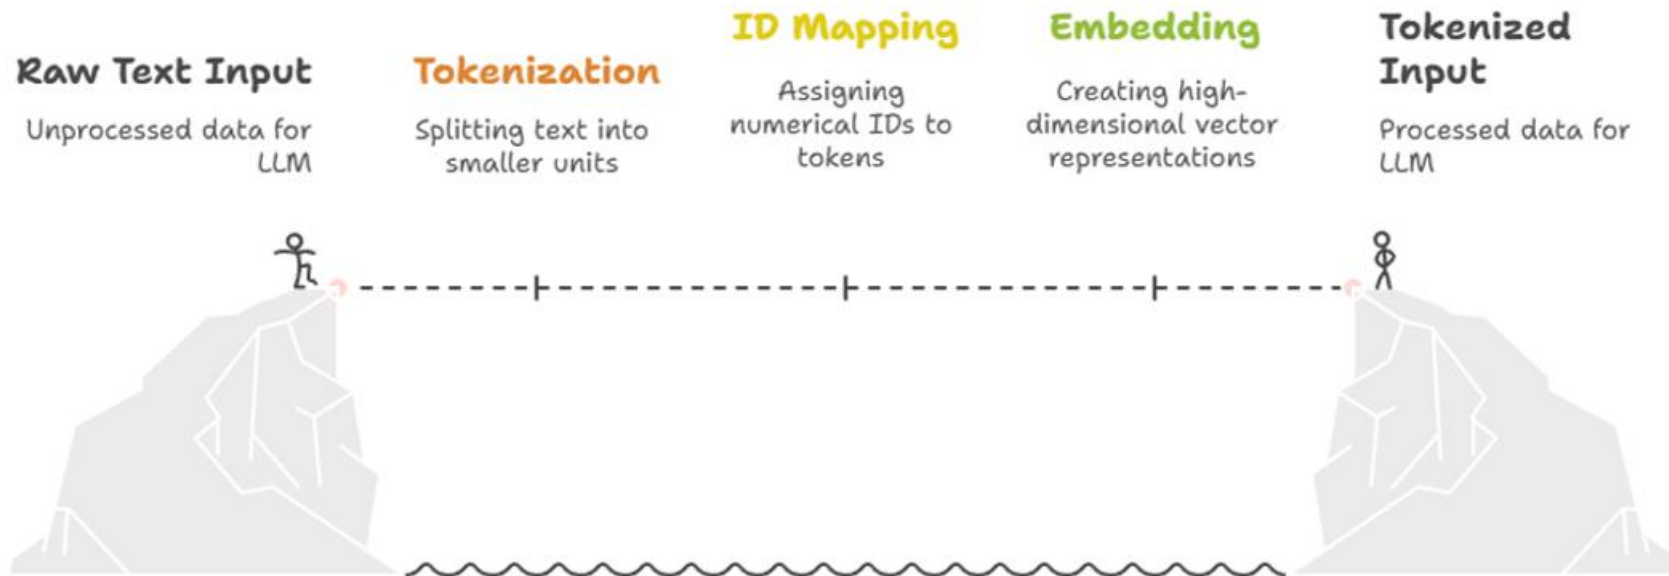


Tokens: The building blocks of LLMs

- **What's a token:** The smallest chunk of text an LLM understands. It could be a whole word, part-word, punctuation, or even a single character.
- **Tokenization = step #1:**
 - Every prompt and training document is split into tokens, forming the model's vocabulary.
 - Example: I heard a dog bark loudly at a cat → 9 tokens.
- **Three common tokenization methods:** Word, character, sub-word. Each trades flexibility vs. efficiency.

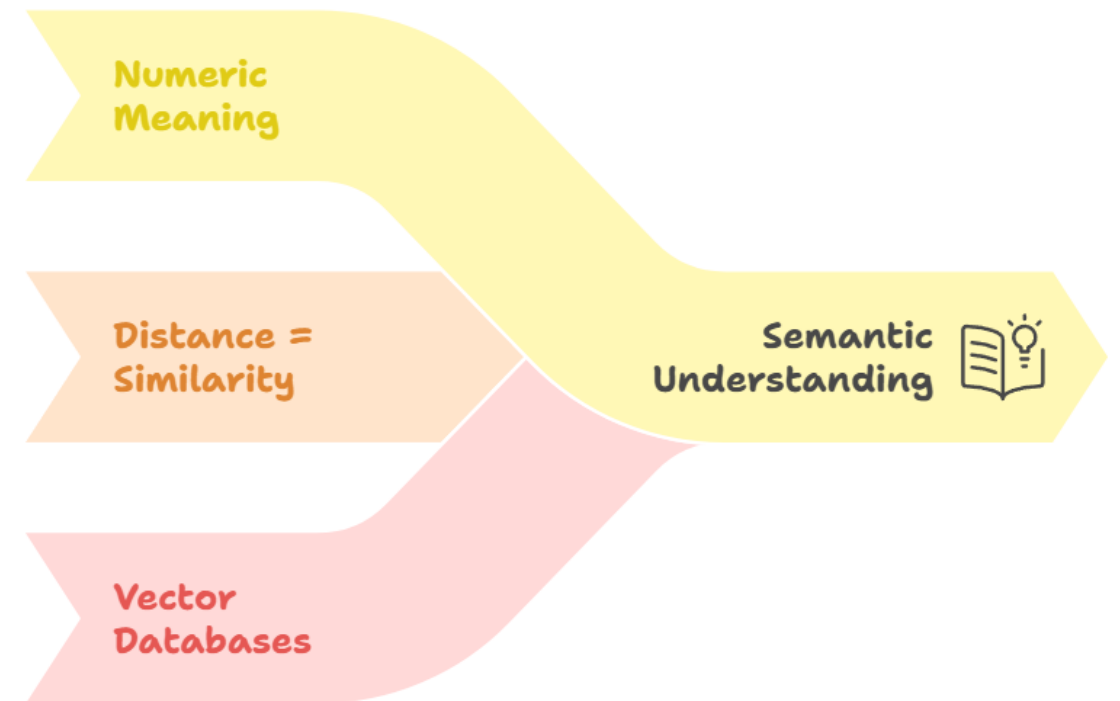
Tokens: The building blocks of LLMs

- **Token IDs & embeddings:** Tokens are mapped to numeric IDs, then to high-dimensional embeddings that capture meaning and context.
- **Context window & limits:** Every model has a max number of input + output tokens (the context window). Go over and your request is clipped or rejected.
- **Why developers should care:** Latency, cost, and rate limits are all measured in tokens, not characters. Concise prompts save money and time.



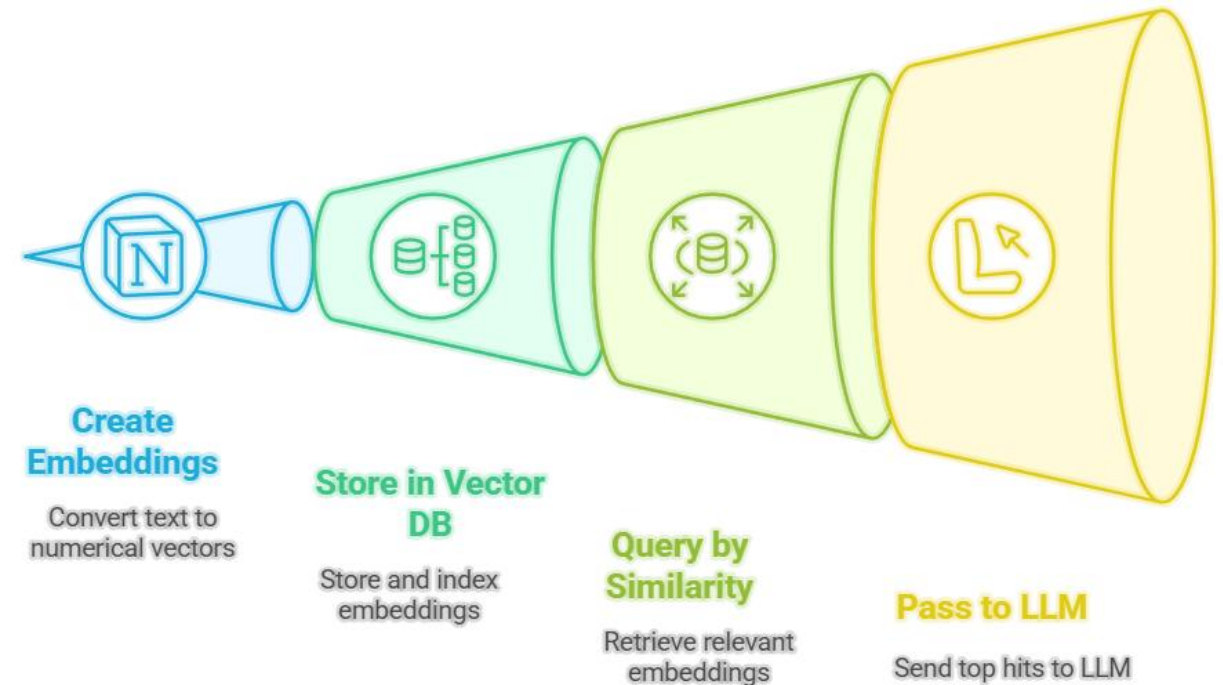
Embeddings: Giving AI a sense of meaning

- **Numeric meaning:** An embedding is a dense vector that captures the semantic essence of text, code, audio, or images.
- **Distance = similarity:** Vectors that point in a similar direction mean their originals are concept-wise close. Cosine similarity is the usual yard-stick.
- **Where they live:** Store millions of vectors in a *vector database*, like Azure AI Search, Cosmos DB, Postgres, etc.
- **Why developers should care:**
 - Retrieval-augmented generation (RAG) with your data.
 - Summaries & prompt compression to fit token limits
 - Classification, translation, recommendation, multimodal mash-ups, etc.




Vector Databases: Long-term memory for your .NET AI apps

- **What they are:** Databases built to *store & index* high-dimensional **embeddings** so you can search by *meaning* instead of keywords.
- **Why we use them:** Similar-item discovery, product or content recommendations, anomaly & fraud detection, and **Retrieval-Augmented Generation (RAG)** for LLM chat with your own data.
- **Typical vector search workflow:**
 - Create embeddings with Azure OpenAI.
 - Store / index in a vector DB (Azure AI Search, Cosmos DB, etc.).
 - Convert user query to an embedding, then run a **nearest-neighbor search** (cosine/Euclidean) inside vector database and query by similarity.
 - Pass top hits to LLM.



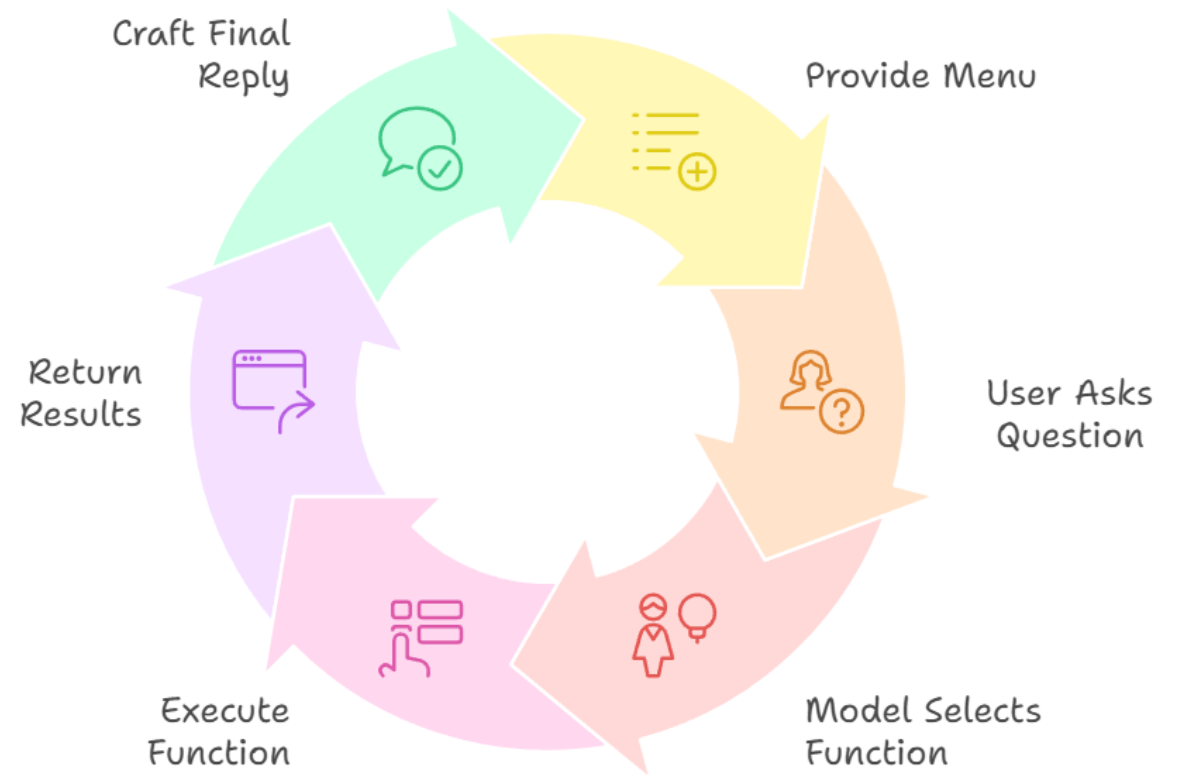
Retrieval-Augmented Generation (RAG)

- Brings your own data into the conversation without re-training the LLM.
- Workflow: Chunk → Embed → Store → Retrieve → Prompt.
- Uses vector search to find the most relevant pieces of content.
- Cheaper & faster than fine-tuning. Keeps answers fresh and source-linked.
- Easily wired into .NET with Semantic Kernel + Azure AI Search / Cosmos DB.



OpenAI Function Calling

- **Describe each tool as JSON:** The model returns which function to run + arguments instead of pure text.
- **Turns the LLM into a router:** For real-time APIs, databases, or code. No extra fine-tuning needed.
- **Workflow:** Prompt → JSON call(s) → execute → feed results → final answer.
- Built into Semantic Kernel with **[KernelFunction]** attributes on C# methods



Semantic Kernel's "Kernel" (The mission-control center)

- **Central DI container:** Holds every AI service & plugin your app needs.
- **Orchestrates the whole cycle:** Pick service → build prompt → send → parse → return.
- **Key Semantic Kernel components**
 - AI service connectors
 - Vector-store (aka. memory) connectors
 - Functions & Plugins
 - Prompt templates
 - Filters

```
builder.Services
    .AddKernel()
    .AddAzureOpenAIChatCompletion(builder.Configuration["AzureDeployment"]!)
    .AddAzureAISearchVectorStore()
    .AddAzureOpenAITextEmbeddingGeneration(builder.Configuration["EmbeddingModelDeployment"]!)
    .ConfigureOpenTelemetry(builder.Configuration);
```

Semantic Kernel Agent Framework

- Adds autonomous, goal-driven agents on top of familiar Semantic Kernel patterns.
- Base **Agent** class: Concrete types like **ChatCompletionAgent**, **AzureAIAgent**, **OpenAIAssistantAgent** (all kernel-powered).
- Agents can co-operate in an **AgentChat** / **AgentGroupChat** for multi-step or multi-agent workflows.
- Each agent links to a Kernel + Plugins / Function Calling: Real tools & memory inside the loop.

```
2 references | Orestis Meikopoulos, 6 days ago | 1 author, 2 changes
public static ChatCompletionAgent CreateChatCompletionAgent(
    Kernel kernel,
    AgentType agentType)
{
    var agentPromptTemplateConfig = SystemPromptFactory
        .GetAgentPromptTemplateConfig(agentType);

    return new ChatCompletionAgent(
        agentPromptTemplateConfig,
        new LiquidPromptTemplateFactory()
    {
        Name = SystemPromptFactory.GetAgentName(agentType),
        Description = agentPromptTemplateConfig.Description,
        Instructions = $""{agentPromptTemplateConfig.Template}""",
        Kernel = PluginFactory.GetAgentKernel(
            kernel,
            agentType,
            kernel.LoggerFactory),
        Arguments = CreateFunctionChoiceAutoBehavior(),
        LoggerFactory = kernel.LoggerFactory
    });
}
```

Azure AI Agent Service: Managed, server-side Agents

- **Fully managed:** Build, deploy, and scale AI agents without running any infrastructure yourself.
- **Server-side tool calling & state:** Azure handles the entire JSON-call → invoke → response loop and persists conversation threads.
- **Rich built-in tools:** File Search, Code Interpreter, Bing, Azure AI Search, Azure Functions & bring your own.
- **Model choice freedom:** Mix Azure OpenAI, Llama 3, Mistral, under the same API.
- **Enterprise extras:** Keyless auth, BYO or managed storage, Responsible-AI filters, elastic scaling.

Demo








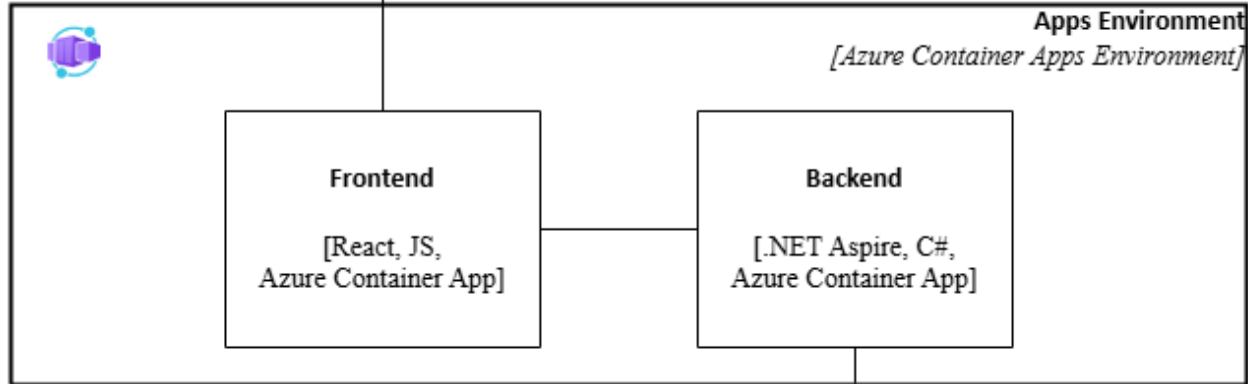




Users

Coding Influencer System

[Azure]





Thank you

- For the opportunity
- For participating
- For listening